

Problem description

You are to predict whether a company will go bankrupt in the following year, based on financial attributes of the company.

Perhaps you are contemplating lending money to a company, and need to know whether the company is in near-term danger of not being able to repay.

Import modules

```
In [7]: ► ## Standard imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import sklearn

import os
import math

%matplotlib inline
```

API

```
In [8]: ► ## Load the bankruptcy_helper module

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

# Reload all modules imported with %aimport
%load_ext autoreload
%autoreload 1

# Import bankruptcy_helper module
import bankruptcy_helper
%aimport bankruptcy_helper

helper = bankruptcy_helper.Helper()
```

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

Get the data

The first step in our Recipe is Get the Data.

- Each example is a row of data corresponding to a single company
- There are 64 attributes, described in the section below
- The column `Bankrupt` is 1 if the company subsequently went bankrupt; 0 if it did not go bankrupt
- The column `Id` is a Company Identifier

```
In [9]: ▶ # Data directory
DATA_DIR = "./Data"

if not os.path.isdir(DATA_DIR):
    DATA_DIR = "../resource/asnlib/publicdata/bankruptcy/data"

data_file = "5th_yr.csv"
data = pd.read_csv( os.path.join(DATA_DIR, "train", data_file) )

target_attr = "Bankrupt"

n_samples, n_attrs = data.shape
print("Data shape: ", data.shape)

Data shape: (4818, 66)
```

Have a look at the data

```
In [10]: ▶ data.head()
```

Out[10]:

	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10	
0	0.025417	0.41769	0.0568	1.1605	-126.39	0.41355	0.025417	1.2395	1.16500	0.51773	0.0
1	-0.023834	0.2101	0.50839	4.2374	22.034	0.058412	-0.027621	3.6579	0.98183	0.76855	-0.0
2	0.030515	0.44606	0.19569	1.565	35.766	0.28196	0.039264	0.88456	1.05260	0.39457	0.0
3	0.052318	0.056366	0.54562	10.68	438.2	0.13649	0.058164	10.853	1.02790	0.61173	0.0
4	0.000992	0.49712	0.12316	1.3036	-71.398	0	0.001007	1.0116	1.29210	0.50288	0.0

In [23]: ▶ `print(data.head())`

	X1	X2	X3	X4	X5	X6	X7	X8
X9	X10	X11	X12	X13	X14	X15	X16	X17
X18	X19	X20 \						
0	0.025417	0.417690	0.05680	1.1605	-126.390	0.413550	0.025417	1.23950
1	1.16500	0.51773	0.025417	0.071822	0.042589	0.025417	4553.60	0.080156
2	3941	0.025417	0.032332	64.985				
3	1	-0.023834	0.210100	0.50839	4.2374	22.034	0.058412	-0.027621
4	0.98183	0.76855	-0.027621	-0.175880	0.011274	-0.027621	3138.40	0.116300
5	7595	-0.027621	-0.012744	62.936				
6	2	0.030515	0.446060	0.19569	1.5650	35.766	0.281960	0.039264
7	1.05260	0.39457	0.039264	0.113360	0.056198	0.039264	2400.00	0.152090
8	2	2418	0.039264	0.032526	24.253			
9	3	0.052318	0.056366	0.54562	10.6800	438.200	0.136490	0.058164
10	1.02790	0.61173	0.058164	1.031900	0.186870	0.058164	260.91	1.399000
11	17	7410	0.058164	0.137840	60.675			
12	4	0.000992	0.497120	0.12316	1.3036	-71.398	0.000000	0.001007
13	1.01160	1.29210	0.50288	0.017035	0.002484	0.040636	0.001007	3455.90
14	2	0116	0.001007	0.000780	104.830			

	X21	X22	X23	X24	X25	X26	X27	X28
X29	X30	X31	X32	X33	X34	X35	X36	X
37	X38	X39	X40 \					
0	1.09350	0.113860	0.032332	0.413550	0.51773	0.080156	1.687400	0.096384
1	4.8078	0.513970	0.032332	191.430	1.9067	0.272590	0.113860	0.79532
2	4.2	43700	0.58152	0.144830	0.464990			
3	1	1.15070	-0.015089	-0.010996	0.077482	0.76855	0.134330	-0.068354
4	1.519500	4.8829	0.092599	-0.012744	25.965	14.0570	-0.071818	-0.015089
5	5.4	97200	0.82161	-0.006962	0.077370			
6	2	0.82243	0.038380	0.025278	0.353230	0.39457	0.132470	0.334650
7	0.427330	4.5257	0.345420	0.032526	110.230	3.3111	0.086042	0.038380
8	4.6	32700	0.49426	0.031794	0.096624			
9	3	0.94921	0.059565	0.123980	0.149690	0.61173	1.295300	1.451000
10	1.370900	4.7352	-0.990460	0.137840	50.118	7.2828	1.056700	0.059565
11	100.1	85857	0.61173	0.141160	8.432300			
12	4	0.97245	0.017035	0.000768	0.002107	0.10133	0.105590	1.062900
13	0.261370	4.0033	0.381750	0.012043	112.330	3.2495	2.651300	-0.025970
14	15.8	90000	0.51280	-0.020100	0.017617			

	X41	X42	X43	X44	X45	X46	X47	X48
X49	X50	X51	X52	X53	X54	X55	X56	X57
X58	X59	X60 \						
0	0.11263	0.144830	114.28	49.297	0.181600	0.76501	75.710	0.105790
1	0.1	34580	0.98325	0.353890	0.524470	0.87853	0.98679	3649.0
2	0.0490	94	0.85835	0.123220	5.6167			
3	1	0.18686	-0.006962	110.01	47.079	-0.063774	1.85760	61.792
4	-0.067146	-0.0	30980	3.16710	0.157040	0.071138	2.29710	2.45570
5	38823.0	-0.018502	-0.0310	11	1.01850	0.069047	5.7996	
6	2	0.21903	0.031794	153.78	129.530	0.380420	1.33340	25.528
7	0.009805	0.0	08122	1.21520	0.346370	0.302010	0.86161	1.07930
8	0.0773	37	0.95006	0.252660	15.0490			
9	3	0.02309	0.141160	109.59	48.911	0.745850	9.43550	62.370
10	0.038873	0.0	92122	10.68000	0.056366	0.137310	1.53700	1.53700
11	29652.0	0.027178	0.0855	24	0.97282	0.000000	6.0157	
12	4	0.24180	0.013184	147.36	42.531	0.002675	0.38879	102.760
13	-0.034462	-0.0	26672	1.06370	0.405620	0.307740	1.06720	1.08820
14	0.0019	74	0.99925	0.019736	3.4819			

	X61	X62	X63	X64	Bankrupt	Id
0	7.4042	164.310	2.2214	1.3340	0	4510
1	7.7529	26.446	13.8020	6.4782	0	3537
2	2.8179	104.730	3.4852	2.6361	0	3920
3	7.4626	48.756	7.4863	1.0602	0	1806
4	8.5820	114.580	3.1854	2.7420	0	1529

Pretty *unhelpful* !

What are these mysteriously named features ?

Description of attributes

Attribute Information:

Id Company Identifier

- X1 net profit / total assets
- X2 total liabilities / total assets
- X3 working capital / total assets
- X4 current assets / short-term liabilities
- X5 [(cash + short-term securities + receivables - short-term liabilities) / (operating expenses - depreciation)] * 365
- X6 retained earnings / total assets
- X7 EBIT / total assets
- X8 book value of equity / total liabilities
- X9 sales / total assets
- X10 equity / total assets
- X11 (gross profit + extraordinary items + financial expenses) / total assets
- X12 gross profit / short-term liabilities
- X13 (gross profit + depreciation) / sales
- X14 (gross profit + interest) / total assets
- X15 (total liabilities * 365) / (gross profit + depreciation)
- X16 (gross profit + depreciation) / total liabilities
- X17 total assets / total liabilities
- X18 gross profit / total assets
- X19 gross profit / sales
- X20 (inventory * 365) / sales
- X21 sales (n) / sales (n-1)
- X22 profit on operating activities / total assets
- X23 net profit / sales
- X24 gross profit (in 3 years) / total assets
- X25 (equity - share capital) / total assets
- X26 (net profit + depreciation) / total liabilities
- X27 profit on operating activities / financial expenses
- X28 working capital / fixed assets
- X29 logarithm of total assets
- X30 (total liabilities - cash) / sales
- X31 (gross profit + interest) / sales
- X32 (current liabilities * 365) / cost of products sold
- X33 operating expenses / short-term liabilities
- X34 operating expenses / total liabilities
- X35 profit on sales / total assets
- X36 total sales / total assets
- X37 (current assets - inventories) / long-term liabilities
- X38 constant capital / total assets
- X39 profit on sales / sales
- X40 (current assets - inventory - receivables) / short-term liabilities
- X41 total liabilities / ((profit on operating activities + depreciation) * (12/365))
- X42 profit on operating activities / sales
- X43 rotation receivables + inventory turnover in days
- X44 (receivables * 365) / sales
- X45 net profit / inventory
- X46 (current assets - inventory) / short-term liabilities
- X47 (inventory * 365) / cost of products sold
- X48 EBITDA (profit on operating activities - depreciation) / total assets
- X49 EBITDA (profit on operating activities - depreciation) / sales
- X50 current assets / total liabilities
- X51 short-term liabilities / total assets
- X52 (short-term liabilities * 365) / cost of products sold)

- X53 equity / fixed assets
- X54 constant capital / fixed assets
- X55 working capital
- X56 (sales - cost of products sold) / sales
- X57 (current assets - inventory - short-term liabilities) / (sales - gross profit - depreciation)
- X58 total costs / total sales
- X59 long-term liabilities / equity
- X60 sales / inventory
- X61 sales / receivables
- X62 (short-term liabilities * 365) / sales
- X63 sales / short-term liabilities
- X64 sales / fixed assets

This may still be somewhat unhelpful for those of you not used to reading Financial Statements.

But that's partially the point of the exercise

- You can *still* perform Machine Learning *even if* you are not an expert in the problem domain
 - That's what makes this a good interview exercise: you can demonstrate your thought process even if you don't know the exact meaning of the terms
- Of course: becoming an expert in the domain *will improve* your ability to create better models
 - Feature engineering is easier if you understand the features, their inter-relationships, and the relationship to the target

Let's get a feel for the data

- What is the type of each attribute ?

In [11]: ▶ data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4818 entries, 0 to 4817
Data columns (total 66 columns):
#   Column      Non-Null Count  Dtype
---  -
0   X1           4818 non-null   object
1   X2           4818 non-null   object
2   X3           4818 non-null   object
3   X4           4818 non-null   object
4   X5           4818 non-null   object
5   X6           4818 non-null   object
6   X7           4818 non-null   object
7   X8           4818 non-null   object
8   X9           4818 non-null   float64
9   X10          4818 non-null   object
10  X11          4818 non-null   object
11  X12          4818 non-null   object
12  X13          4818 non-null   float64
13  X14          4818 non-null   object
14  X15          4818 non-null   object
15  X16          4818 non-null   object
16  X17          4818 non-null   object
17  X18          4818 non-null   object
18  X19          4818 non-null   float64
19  X20          4818 non-null   float64
20  X21          4818 non-null   object
21  X22          4818 non-null   object
22  X23          4818 non-null   float64
23  X24          4818 non-null   object
24  X25          4818 non-null   object
25  X26          4818 non-null   object
26  X27          4818 non-null   object
27  X28          4818 non-null   object
28  X29          4818 non-null   object
29  X30          4818 non-null   float64
30  X31          4818 non-null   float64
31  X32          4818 non-null   object
32  X33          4818 non-null   object
33  X34          4818 non-null   object
34  X35          4818 non-null   object
35  X36          4818 non-null   object
36  X37          4818 non-null   object
37  X38          4818 non-null   object
38  X39          4818 non-null   float64
39  X40          4818 non-null   object
40  X41          4818 non-null   object
41  X42          4818 non-null   float64
42  X43          4818 non-null   float64
43  X44          4818 non-null   float64
44  X45          4818 non-null   object
45  X46          4818 non-null   object
46  X47          4818 non-null   object
47  X48          4818 non-null   object
48  X49          4818 non-null   float64
49  X50          4818 non-null   object
50  X51          4818 non-null   object
51  X52          4818 non-null   object
52  X53          4818 non-null   object
53  X54          4818 non-null   object
54  X55          4818 non-null   float64
55  X56          4818 non-null   float64
56  X57          4818 non-null   object
57  X58          4818 non-null   float64
58  X59          4818 non-null   object
59  X60          4818 non-null   object

```



```

60 X61          4818 non-null    object
61 X62          4818 non-null    float64
62 X63          4818 non-null    object
63 X64          4818 non-null    object
64 Bankrupt     4818 non-null    int64
65 Id           4818 non-null    int64
dtypes: float64(16), int64(2), object(48)
memory usage: 2.4+ MB

```

You may be puzzled:

- Most attributes are `object` and *not* numeric (`float64`)
- But looking at the data via `data.head()` certainly gives the impression that all attributes are numeric

Welcome to the world of messy data ! The dataset has represented numbers as strings.

- These little unexpected challenges are common in the real-world
- Data is rarely perfect and clean

So you might want to first convert all attributes to numeric

Hint

- Look up the Pandas method `to_numeric`
 - We suggest you use the option `errors='coerce'`

Evaluating your project

```

In [12]: ▶ holdout_data = pd.read_csv( os.path.join(DATA_DIR, "holdout", '5th_yr.csv') )

print("Data shape: ", holdout_data.shape)

```

```
Data shape: (1092, 65)
```

We will evaluate your model on the holdout examples using metrics

- Accuracy
- Recall
- Precision

From our lecture: we may have to make a trade-off between Recall and Precision.

Our evaluation of your submission will be partially based on how you made (and described) the trade-off.

You may assume that it is 5 times worse to *fail to identify a company that will go bankrupt* than it is to fail to identify a company that won't go bankrupt.

Submission guidelines

Although your notebook may contain many models (e.g., due to your iterative development) we will only evaluate a single model. So choose one (explain why !) and do the following.

- You will implement the body of a subroutine `MyModel`
 - That takes as argument a Pandas DataFrame
 - Each row is an example on which to predict

- The features of the example are elements of the row
- Performs predictions on each example
- Returns an array of predictions with a one-to-one correspondence with the examples in the test set

We will evaluate your model against the holdout data

- By reading the holdout examples `X_hold` (as above)
- Calling `y_hold_pred = MyModel(X_hold)` to get the predictions
- Comparing the predicted values `y_hold_pred` against the true labels `y_hold` which are known only to the instructors

See the following cell as an illustration:

```
X_hold = pd.read_csv( os.path.join(DATA_DIR, "holdout", '5th_yr.csv') )
```

Predict using MyModel

```
y_hold_pred = MyModel(X_hold)
```

Compute metrics

accuracy

```
accuracy_hold = accuracy_score(y_hold, y_hold_pred)
```

recall_

```
recall_hold = recall_score(y_hold, y_hold_pred, pos_label=1, average="binary")
```

precision

```
precision_hold = precision_score(y_hold, y_hold_pred, pos_label=1, average="binary")
```

```
print("\t{m:s} Accuracy: {a:3.1%}, Recall {r:3.1%}, Precision {p:3.1%}".format(m=name, a=accuracy_hold, r=recall_hold, p=precision_hold ) )
```

```
In [13]: ▶ from sklearn.decomposition import PCA
from sklearn import svm
from sklearn.metrics import accuracy_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_predict, train_test_split
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn import tree
import math
```

```
In [22]: # Convert all data values to numeric, coercing errors to NaN
data = data.applymap(pd.to_numeric, errors='coerce')
holdout_data = holdout_data.applymap(pd.to_numeric, errors='coerce')

# Replace '?' with NaN to handle missing values
data.replace('?', np.nan, inplace=True)
data.fillna(data.mean(), inplace=True)

# Fill missing values with column-wise mean
holdout_data.replace('?', np.nan, inplace=True)
holdout_data.fillna(holdout_data.mean(), inplace=True)

# Convert DataFrame to NumPy array
df = data.values
out = holdout_data.values

# Display summary statistics of the dataset

data.describe()

# Extract features (all columns except last two) and normalize them
X=df[:, :-2]
X = StandardScaler().fit_transform(X) # Standardize features for ML models

# Extract target variable (second last column)
y=df[:, -2]
# Alternative feature and target extraction using Pandas indexing
X_data=data.iloc[:, :-2]
y_data=data.iloc[:, -2]
```

Out[22]:

	X1	X2	X3	X4	X5	X6	X7
count	4818.000000	4818.000000	4818.000000	4818.000000	4.818000e+03	4818.000000	4818.000000
mean	-0.055232	0.533272	0.188992	4.978602	1.956441e+01	-0.070060	-0.042516
std	6.705958	1.202660	1.282164	100.117705	2.382015e+04	7.776908	6.706577
min	-463.890000	0.000000	-72.067000	0.000000	-1.076400e+06	-463.890000	-463.890000
25%	0.004042	0.254765	0.044978	1.101525	-4.290075e+01	0.000000	0.005977
50%	0.046428	0.451610	0.218155	1.645450	4.932050e-01	0.000000	0.056653
75%	0.116725	0.662140	0.420033	2.943000	4.978025e+01	0.110387	0.135972
max	2.352300	72.416000	28.336000	6845.800000	1.250100e+06	203.150000	2.352300

```
In [21]: ▶ pd.set_option('display.max_columns', None)
pd.set_option('display.expand_frame_repr', True)
pd.set_option('display.width', 200)
print(data.describe())
```

	X1	X2	X3	X4	X5	
X6	X7	X8	X9	X10	X11	X12
X13	X14 \					
count	4818.000000	4818.000000	4818.000000	4818.000000	4.818000e+03	4818.00
0000	4818.000000	4818.000000	4818.000000	4818.000000	4818.000000	4818.0000
00	4818.000000	4818.000000				
mean	-0.055232	0.533272	0.188992	4.978602	1.956441e+01	-0.07
0060	-0.042516	5.739985	1.579277	0.503606	-0.024637	1.2183
57	0.452284	-0.042485				
std	6.705958	1.202660	1.282164	100.117705	2.382015e+04	7.77
6908	6.706577	109.348749	1.342723	4.043196	6.706133	38.7865
94	34.196231	6.706578				
min	-463.890000	0.000000	-72.067000	0.000000	-1.076400e+06	-463.89
0000	-463.890000	-3.735100	0.000191	-71.444000	-463.890000	-96.2390
00	-310.340000	-463.890000				
25%	0.004042	0.254765	0.044978	1.101525	-4.290075e+01	0.00
0000	0.005977	0.482350	1.015600	0.319095	0.016183	0.0178
81	0.024954	0.005977				
50%	0.046428	0.451610	0.218155	1.645450	4.932050e-01	0.00
0000	0.056653	1.154350	1.140500	0.522195	0.071071	0.1686
85	0.067723	0.056685				
75%	0.116725	0.662140	0.420033	2.943000	4.978025e+01	0.11
0387	0.135972	2.814600	1.814050	0.721670	0.147890	0.5470
95	0.134847	0.136055				
max	2.352300	72.416000	28.336000	6845.800000	1.250100e+06	203.15
0000	2.352300	6868.500000	37.807000	266.860000	2.352300	2470.3000
00	2340.200000	2.352300				

	X15	X16	X17	X18	X19	
X20	X21	X22	X23	X24	X25	X2
6	X27	X28 \				
count	4.818000e+03	4818.000000	4818.000000	4818.000000	4818.000000	4818.00
0000	4818.000000	4818.000000	4818.000000	4818.000000	4818.000000	4818.0000
00	4818.000000	4818.000000				
mean	2.768033e+03	1.261838	6.840618	-0.031071	-0.082028	57.00
2168	2.740360	-0.015683	-0.090357	0.137708	0.371927	1.1621
32	381.584380	6.822847				
std	3.561790e+04	40.952597	109.355403	6.753507	5.754879	182.01
8911	110.389839	6.228986	5.725258	8.040588	4.109756	38.8129
95	9862.743624	108.277557				
min	-4.667700e+05	-52.440000	0.000000	-463.890000	-310.800000	-29.34
0000	-135.150000	-431.590000	-310.890000	-463.890000	-71.444000	-52.4540
00	-158130.000000	-1089.700000				
25%	2.492100e+02	0.076804	1.509850	0.005977	0.004368	18.53
7750	0.992880	0.000668	0.002728	0.026527	0.180500	0.0703
24	0.145345	0.091070				
50%	8.972850e+02	0.236445	2.214850	0.056685	0.035307	38.62
3000	1.122300	0.062017	0.030074	0.154970	0.425990	0.2108
80	1.162200	0.543290				
75%	2.311825e+03	0.638502	3.926975	0.136055	0.088287	66.85
0750	1.274500	0.136815	0.075800	0.350295	0.638653	0.5809
82	6.267475	1.732375				
max	1.341700e+06	2837.400000	6869.500000	55.125000	77.244000	9928.50
0000	7661.500000	15.541000	77.244000	252.340000	266.860000	2689.1000
00	565940.000000	5534.000000				

	X29	X30	X31	X32	X33	
X34	X35	X36	X37	X38	X39	X
40	X41	X42 \				
count	4818.000000	4818.000000	4818.000000	4.818000e+03	4818.000000	4818.00
0000	4818.000000	4818.000000	4818.000000	4818.000000	4818.000000	4818.0000
000	4818.000000	4818.000000				
mean	4.161201	0.691991	-0.068612	2.074823e+03	8.520843	5.07
4955	-0.017852	2.075176	100.185857	0.589316	0.019158	2.225
580	2.556513	-0.014902				

std	0.829921	11.217691	5.748967	8.728196e+04	112.335444	109.97
3338	6.227517	14.101975	899.774590	4.035992	0.756891	63.131
677	91.177510	2.617310				
min	0.006359	-23.060000	-310.800000	0.000000e+00	0.000000	-16.01
5000	-431.590000	0.000191	-4.325800	-71.444000	-47.047000	-0.335
160	-269.990000	-143.520000				
25%	3.656250	0.085560	0.007006	5.013275e+01	2.782500	0.24
9400	0.008295	1.063300	2.846125	0.433462	0.005536	0.053
029	0.031838	0.000687				
50%	4.175200	0.225955	0.042826	8.197500e+01	4.488750	1.71
4700	0.061804	1.552850	56.770000	0.620230	0.040252	0.183
200	0.094913	0.040963				
75%	4.669700	0.406145	0.098331	1.321750e+02	7.298550	3.97
7275	0.138695	2.232275	100.185857	0.774872	0.091854	0.674
010	0.217020	0.090831				
max	9.698300	656.450000	77.244000	4.277200e+06	7590.500000	7590.50
0000	15.541000	965.660000	40559.000000	266.860000	2.901100	4303.200
000	5043.300000	40.386000				

	X43	X44	X45	X46	X47	
X48	X49	X50	X51	X52	X53	X5
4	X55	X56 \				
count	4818.000000	4818.000000	4818.000000	4818.000000	4818.000000	48
18.000000	4818.000000	4818.000000	4818.000000	4818.000000	4818.000000	481
8.000000	4.818000e+03	4818.000000				
mean	155.612840	98.610765	81.795906	4.103433	153.579739	
-0.103035	-0.072253	4.192413	0.426457	0.817538	8.089876	
9.142221	1.074023e+04	0.056109				
std	795.989622	725.594072	5273.616963	100.015679	3617.994648	
7.829928	2.638443	99.598952	1.180774	14.189417	144.828893	160.
361029	8.284685e+04	0.755462				
min	-3975.600000	-3946.200000	-3037.300000	-0.108060	-18.658000	-5
42.560000	-144.800000	0.000000	0.000000	0.000000	-1088.700000	-108
8.700000	-1.118500e+06	-46.788000				
25%	76.214500	39.050250	0.026996	0.646130	19.849250	
-0.030394	-0.022447	0.834400	0.186087	0.137313	0.764355	
1.001350	9.771450e+01	0.011478				
50%	106.670000	59.014000	0.291190	1.066600	42.389500	
0.019769	0.012481	1.289150	0.328945	0.223895	1.296650	1.
440100	1.829500e+03	0.053663				
75%	149.365000	86.087750	0.943597	2.000750	74.090750	
0.095739	0.060499	2.372050	0.519905	0.360503	2.439775	2.
576300	7.786950e+03	0.124030				
max	40515.000000	40515.000000	366030.000000	6845.800000	185610.000000	
15.541000	16.866000	6845.800000	72.416000	666.110000	8309.600000	830
9.600000	4.212200e+06	1.000000				

	X57	X58	X59	X60	X61	
X62	X63	X64	Bankrupt	Id		
count	4818.000000	4818.000000	4818.000000	4.818000e+03	4818.000000	4818.
000000	4818.000000	4818.000000	4818.000000	4818.000000		
mean	0.022793	0.959585	0.273025	1.108795e+03	11.021303	177.
494445	9.287631	38.557533	0.063927	3499.858032		
std	7.247517	0.932427	6.337285	6.942383e+04	43.766529	2279.
713700	113.049493	583.617986	0.244648	1392.049260		
min	-468.670000	-0.085920	-184.980000	-1.244000e+01	-0.092493	0.
000000	0.000000	-3.726500	0.000000	1071.000000		
25%	0.015582	0.876940	0.000000	5.456550e+00	4.236425	45.
065750	3.077000	2.137425	0.000000	2296.250000		
50%	0.108860	0.950825	0.006365	9.449650e+00	6.181450	73.
879500	4.939500	4.224000	0.000000	3500.500000		
75%	0.240200	0.990358	0.208242	1.968150e+01	9.342925	118.
597500	8.097875	9.834275	0.000000	4704.750000		
max	87.981000	47.788000	308.150000	4.818700e+06	1308.500000	127450.
000000	7641.300000	28999.000000	1.000000	5909.000000		

Handling Missing Data:

- Linear Regression Imputation: for data which shows linear trends

Train a linear regression model on the available data
Predict missing values using the trained model
Replace missing values with predicted values

- Interpolation (when missing is between values)

For each missing value:
Find the nearest known data points (x1, y1) and (x2, y2)
Compute slope: $m = (y2 - y1) / (x2 - x1)$
Estimate missing value: $y = y1 + m * (x - x1)$
Replace missing value with estimated y

- Extrapolation (time-series forecasting)

Select the two most recent known data points (x1, y1) and (x2, y2)
Compute slope: $m = (y2 - y1) / (x2 - x1)$
For each missing value:
Predict missing value: $y = y2 + m * (x - x2)$
Replace missing value with estimated y

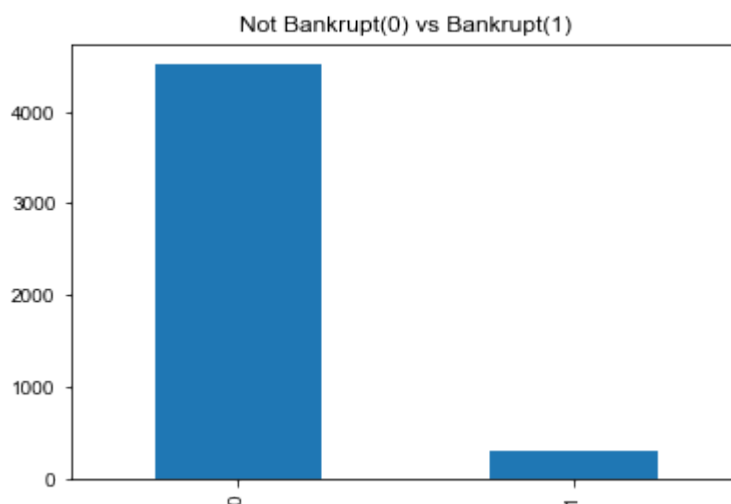
Can be used in ECM to find missing values

Visualize data, and find that the data is imbalanced, and the features have high correlation.

```
In [9]: ▶ import seaborn as sns
y_data.value_counts().plot(kind='bar')
sns.set_style('whitegrid')
sns.set_palette('bwr')
plt.title('Not Bankrupt(0) vs Bankrupt(1)')
plt.show()
```

Out[9]: <AxesSubplot:>

Out[9]: Text(0.5, 1.0, 'Not Bankrupt(0) vs Bankrupt(1)')

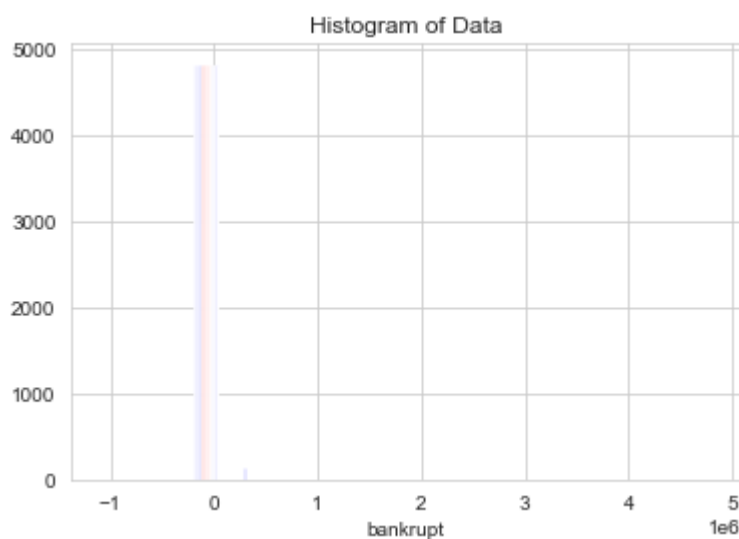


```
In [10]: plt.hist(data, bins=20, alpha=0.5)
plt.xlabel('bankrupt')
plt.title('Histogram of Data')
plt.show()
```

```
Out[10]: (array([[0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                ...,
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.])),
array([-1118500., -821640., -524780., -227920., 68940., 365800.,
        662660., 959520., 1256380., 1553240., 1850100., 2146960.,
        2443820., 2740680., 3037540., 3334400., 3631260., 3928120.,
        4224980., 4521840., 4818700.]),
<a list of 66 BarContainer objects>)
```

```
Out[10]: Text(0.5, 0, 'bankrupt')
```

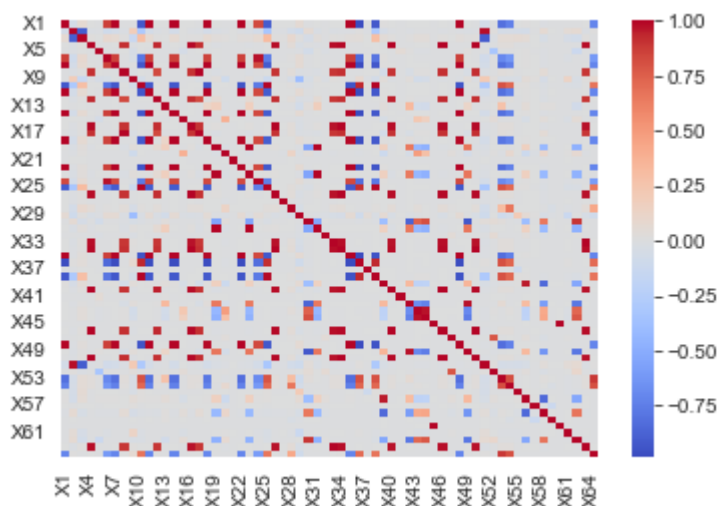
```
Out[10]: Text(0.5, 1.0, 'Histogram of Data')
```



Heatmap

```
In [11]: corr_matrix = X_data.corr()
sns.heatmap(corr_matrix, cmap='coolwarm')
plt.show()
```

```
Out[11]: <AxesSubplot:>
```



Here I use `randomoversampler` to do random oversampling, and use PCA to lower the dimension since we have dataset with features having high correlation with each other.

```
In [12]:  from imblearn.over_sampling import RandomOverSampler
          # Random oversampling
          # ros = RandomOverSampler(random_state=42)
          # X_ros, y_ros = ros.fit_resample(X, y)
```

SMOTE

I also try SMOTE to balance the data to avoid overfitting. After applying SMOTE we can find the recall and precision seems to be balanced.

```
In [13]:  from imblearn.over_sampling import SMOTE

          smote = SMOTE(random_state=42)
          X_ros, y_ros = smote.fit_resample(X, y)
          # X_ros, y_ros = X, y
```

PCA

```
In [14]:  # Deciding the number of components for PCA
          temp = []
          n_components = np.arange(3,30)
          for n in n_components:
              pca = PCA(n_components = n)
              pca.fit_transform(X_ros)
              total_variance = pca.explained_variance_ratio_.sum()
              temp.append(total_variance)
```

```
Out[14]: array([[ -0.34702704,  -0.52941261,  -0.12070452],
                 [ -0.21829935,  -0.92479567,  -0.04201231],
                 [ -0.33985283,  -0.5397412 ,  -0.10783746],
                 ...,
                 [ -0.33001267,  -0.63287133,  -0.10503995],
                 [ -0.64097052,   5.50196878,   0.01719124],
                 [ -0.4842435 ,   2.82453163,  -0.05874565]])
```

```
Out[14]: array([[ -0.34702704,  -0.5294126 ,  -0.12070446,  -0.21890125],
                 [ -0.21829935,  -0.92479567,  -0.04201224,  -0.08630365],
                 [ -0.33985283,  -0.5397412 ,  -0.10783743,  -0.16917515],
                 ...,
                 [ -0.33001267,  -0.63287133,  -0.10503991,  -0.14392445],
                 [ -0.64097052,   5.50196878,   0.01719121,  -0.33927047],
                 [ -0.4842435 ,   2.82453162,  -0.05874572,   0.33723751]])
```

```
Out[14]: array([[ -0.34702704,  -0.5294126 ,  -0.12070438,  -0.21890152,  -0.09077936],
                 [ -0.21829935,  -0.92479567,  -0.04201216,  -0.08630406,  -0.11167949],
                 [ -0.33985283,  -0.5397412 ,  -0.1078374 ,  -0.16917529,  -0.09422986],
                 ...,
                 ...])
```

```
In [15]: sns.set_style('whitegrid')
sns.set_palette('bwr')
plt.figure(figsize=(10,10))
plt.plot(n_components, temp)
plt.xlabel('Number of components')
plt.ylabel('Total explained variance')
plt.title('Total explained variance per components')
plt.show()
```

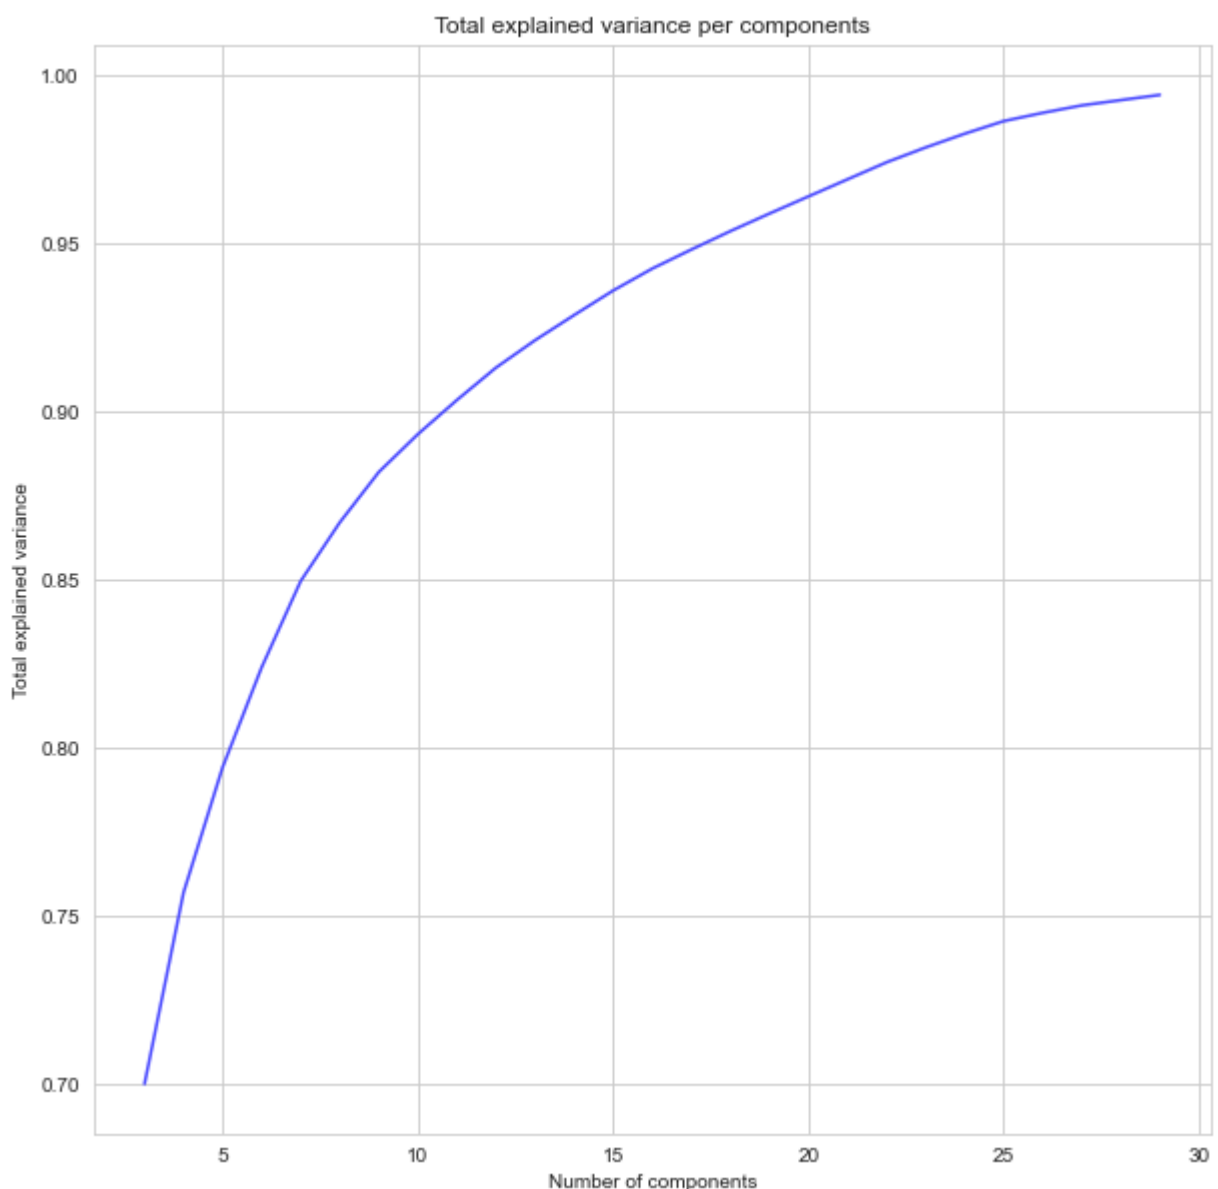
Out[15]: <Figure size 720x720 with 0 Axes>

Out[15]: [<matplotlib.lines.Line2D at 0x16cd9d30fd0>]

Out[15]: Text(0.5, 0, 'Number of components')

Out[15]: Text(0, 0.5, 'Total explained variance')

Out[15]: Text(0.5, 1.0, 'Total explained variance per components')



1. Create an empty list to store explained variance
2. Loop through different numbers of principal components
 - a. Initialize PCA with 'n' components
 - b. Fit PCA to the dataset and transform it
 - c. Calculate the total explained variance
 - d. Store the explained variance in the list
3. Plot the number of components against the total explained variance

In ECM, PCA can be used to identify key risk factors affecting economic capital.

Stress Testing: PCA can identify major risk drivers that contribute to capital adequacy.

Portfolio Risk Management: Helps group correlated financial metrics, improving model robustness.

We can see from the plot that the first 15 components have already explained 92% variance.

```
In [16]: > scaler = StandardScaler()
X_scaled_ros = scaler.fit_transform(X_ros)
pca = PCA(n_components=15)
X_pca_ros = pca.fit_transform(X_scaled_ros)
```

SVM

Since gridsearch takes quite a long time, I use randomsearch instead.

```
In [17]: > from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform
from sklearn import svm
from sklearn.svm import SVC

# Define the hyperparameter search space for SVM
param_dist = {'C': [1, 10, 20], # Regularization Strength, control trade-off
              # between achieving a low error on training
              # data and maintaining model generalization
              'gamma': [0.1, 0.5], # Kernel coefficient for RBF Kernel,
              # control influence of individual training
              # points on decision boundaries
              'kernel': ['linear', 'rbf']} # Choice of Kernel Function, determine
              # how the SVM separates data in
              # feature space

# Initialize an SVM model with a default linear kernel
svm = svm.SVC(kernel='linear')
n_iter_search = 10
random_search = RandomizedSearchCV(svm, param_distributions=param_dist,
                                   n_iter=n_iter_search, cv=5)

# random search helps find the best combination of hyperparameters for a
# machine learning model.
# Fit Randomized Search to the resampled dataset
random_search.fit(X_ros, y_ros)
print("Best hyperparameters: {}".format(random_search.best_params_))
print("Best score: {:.2f}".format(random_search.best_score_))
```

```
Out[17]: RandomizedSearchCV(cv=5, estimator=SVC(kernel='linear'),
                             param_distributions={'C': [1, 10, 20], 'gamma': [0.1, 0.5],
                                                  'kernel': ['linear', 'rbf']})
```

```
Best hyperparameters: {'kernel': 'rbf', 'gamma': 0.5, 'C': 20}
Best score: 0.95
```

```
In [18]: > from sklearn.model_selection import cross_val_score
from sklearn import svm
from sklearn.svm import SVC
print("SVM")

# Initialize an SVM model with specified hyperparameters
svm_model=SVC(C=20, gamma=0.5)
# Split dataset into training (90%) and testing (10%) sets
X_train, X_test, y_train, y_test = train_test_split(X_pca_ros, y_ros,
                                                    test_size=0.1,
                                                    random_state=42)

svm_model.fit(X_train, y_train)
svm_pred = svm_model.predict(X_test)

SVM
```

Out[18]: SVC(C=20, gamma=0.5)

```
In [19]: > name = "SVM model"

# Evaluate model performance
accuracy_test = accuracy_score(y_test, svm_pred)
recall_test = recall_score(y_test, svm_pred, pos_label=1,
                           average="binary")
precision_test = precision_score(y_test, svm_pred, pos_label=1,
                                average="binary", zero_division=1)

print("\t{m:s} Accuracy: {a:3.1%}, Recall {r:3.1%}, Precision {p:3.1%}".format
      (m=name,
       a=accuracy_test,
       r=recall_test,
       p=precision_test)
      )

SVM model Accuracy: 83.6%, Recall 84.1%, Precision 83.0%
```

1. Initialize SVM model with hyperparameters: (depends on cases)
 - C = 20 (higher regularization, reducing misclassification penalties)
 - gamma = 0.5 (higher sensitivity to individual data points)
2. Split dataset into training (90%) and testing (10%) sets.
3. Train SVM model using the training set.
4. Make predictions on the test set.
5. Compute model evaluation metrics:
 - Accuracy: Overall classification correctness
 - Recall: Measures how well bankrupt companies are identified
 - Precision: Measures correctness of bankruptcy predictions
6. Print model performance results.
7. Compute and visualize confusion matrix to analyze prediction errors.

For ECM:

Can be used to identify firms at higher financial risk

Adjust parameters of SVM (c and gamma) to improve accuracy of default probability (Apply RandomizedSearchCV)

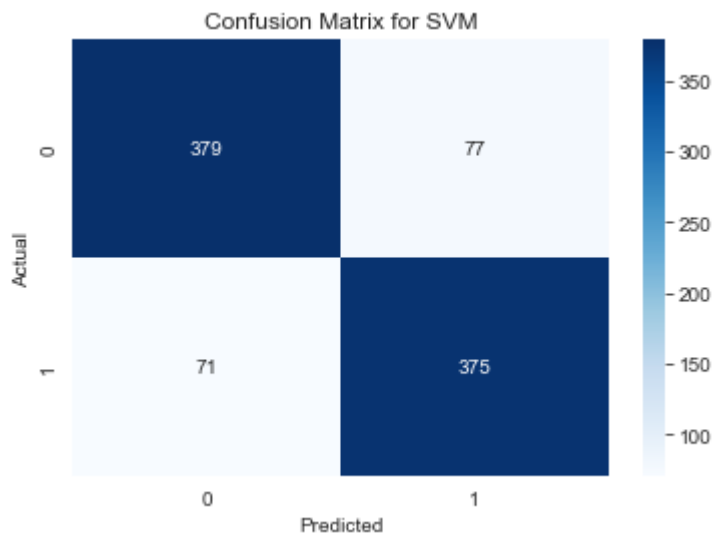
```
In [20]: ► from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, svm_pred)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for SVM')
plt.show()
```

Out[20]: <AxesSubplot:>

Out[20]: Text(0.5, 16.0, 'Predicted')

Out[20]: Text(34.0, 0.5, 'Actual')

Out[20]: Text(0.5, 1.0, 'Confusion Matrix for SVM')



Logistic Regression

```

In [21]: > from sklearn.linear_model import LogisticRegression

# Logistic Regression
print("Logistic Regression")

# Initialize Logistic Regression with balanced class weights to
# handle imbalance
log_reg = LogisticRegression(
    solver='lbfgs', # Solver for optimization
                  # (default, works well for small-medium datasets)
    max_iter=5000, # Increase max iterations to ensure convergence
    class_weight='balanced', # Adjusts weights to balance class distribution
    random_state=42
)

# Train the model on the training set
log_reg.fit(X_train, y_train)
# Predict on the test set
log_reg_pred = log_reg.predict(X_test)

# Evaluate model performance
accuracy_log_reg = accuracy_score(y_test, log_reg_pred)
recall_log_reg = recall_score(y_test, log_reg_pred, pos_label=1,
                             average="binary")
precision_log_reg = precision_score(y_test, log_reg_pred, pos_label=1,
                                    average="binary")

name = "Logistic Regression Model"
print("\t{m:s} Accuracy: {a:3.1%}, Recall {r:3.1%}, Precision {p:3.1%}".format(
    m=name, a=accuracy_log_reg, r=recall_log_reg, p=precision_log_reg
))

# Compute and visualize confusion matrix
cm = confusion_matrix(y_test, log_reg_pred)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Logistic Regression')
plt.show()

```

Logistic Regression

Out[21]: LogisticRegression(class_weight='balanced', max_iter=5000, random_state=42)

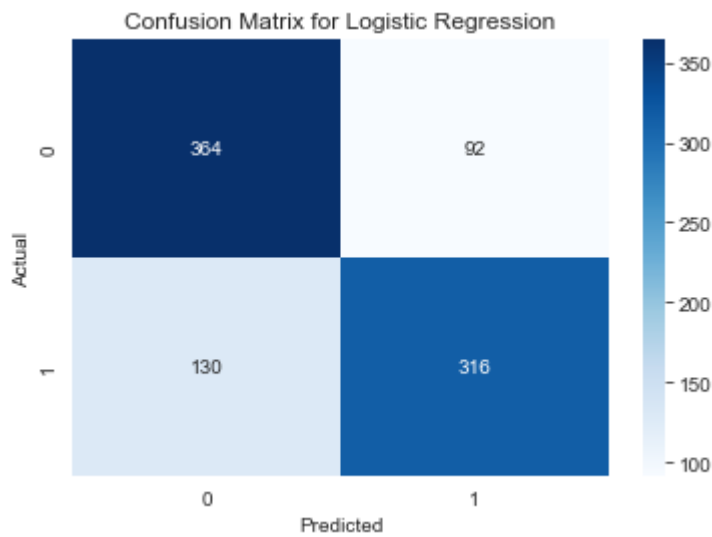
Logistic Regression Model Accuracy: 75.4%, Recall 70.9%, Precision 77.5%

Out[21]: <AxesSubplot:>

Out[21]: Text(0.5, 16.0, 'Predicted')

Out[21]: Text(34.0, 0.5, 'Actual')

Out[21]: Text(0.5, 1.0, 'Confusion Matrix for Logistic Regression')



1. Initialize Logistic Regression with:

- 'lbfgs' solver for optimization
- max_iter=5000 to prevent convergence issues
- class_weight='balanced' to handle imbalanced bankruptcy data

2. Train the model using training data (X_train, y_train).

3. Predict bankruptcy probabilities for test data (X_test).

4. Compute model performance metrics:

- Accuracy: Overall model correctness
- Recall: Ensures bankrupt companies are correctly identified
- Precision: Measures correctness of bankruptcy predictions

5. Print model evaluation results.

6. Compute confusion matrix to analyze classification errors.

7. Visualize confusion matrix using a heatmap.

For ECM:

Interpretable Model: Coefficients represent the impact of financial ratios on bankruptcy probability.

Helps in probability-based economic capital allocation.

Neural network model

```
In [22]: from sklearn.metrics import classification_report
print("Neural network")
print("iteration : 50000")
logistic_activation = MLPClassifier(
    activation = 'logistic', # Uses sigmoid activation function
    hidden_layer_sizes = (90,90,3), # Three-layer neural network structure
    solver = "sgd", # Uses stochastic gradient descent optimizer
    random_state=1, max_iter=50000)

# Train the neural network model
logistic_activation.fit(X_train, y_train)
# Predict on the test set
logis_pred = logistic_activation.predict(X_test)
score_logis = accuracy_score(logis_pred, y_test)
print("sigmoid activation function test accuracy:", score_logis)
```

```
Neural network
iteration : 50000
```

```
Out[22]: MLPClassifier(activation='logistic', hidden_layer_sizes=(90, 90, 3),
                      max_iter=50000, random_state=1, solver='sgd')

sigmoid activation function test accuracy: 0.4911308203991131
```

```
In [23]: name = "Neural network model"

# Evaluate model performance
accuracy_test = accuracy_score(y_test, logis_pred)
recall_test = recall_score(y_test, logis_pred, pos_label=1, average="binary")
precision_test = precision_score(y_test, logis_pred, pos_label=1,
                                average="binary", zero_division=1)

print("\t{m:s} Accuracy: {a:3.1%}, Recall {r:3.1%}, Precision {p:3.1%}".format
      (m=name,
       a=accuracy_test,
       r=recall_test,
       p=precision_test)
      )
```

```
Neural network model Accuracy: 49.1%, Recall 99.3%, Precision 49.3%
```

1. Initialize a Neural Network (MLPClassifier) with:

- Sigmoid activation function ('logistic')
- Three hidden layers: 90, 90, 3 neurons
- Stochastic Gradient Descent (SGD) optimizer
- max_iter=50000 to ensure full convergence

2. Train the model using training data (X_train, y_train).

3. Predict bankruptcy outcomes for test data (X_test).

4. Compute performance metrics:

- Accuracy: Overall correctness of predictions
- Recall: Ability to correctly classify bankrupt companies
- Precision: Accuracy of bankruptcy predictions

5. Print model evaluation results.

For ECM:

Handles nonlinear relationship in financial data

Sigmoid activation (logistic) is useful for probabilistic bankruptcy classification

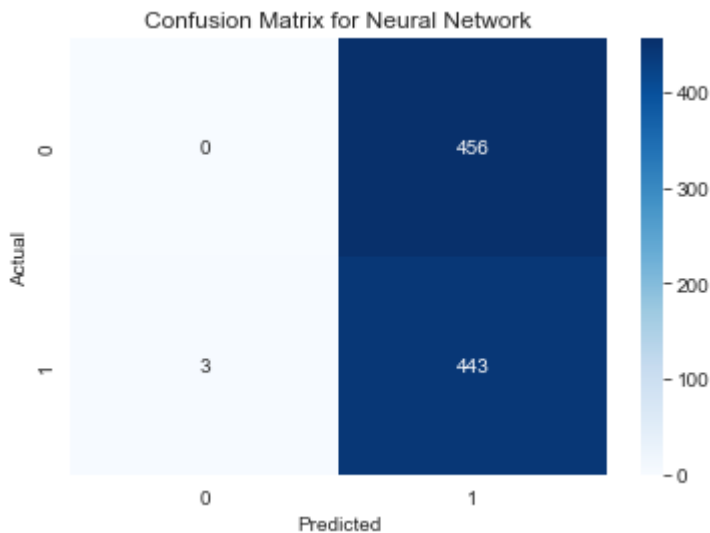

```
In [24]: ► cm = confusion_matrix(y_test, logis_pred)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Neural Network')
plt.show()
```

Out[24]: <AxesSubplot:>

Out[24]: Text(0.5, 16.0, 'Predicted')

Out[24]: Text(34.0, 0.5, 'Actual')

Out[24]: Text(0.5, 1.0, 'Confusion Matrix for Neural Network')



Decision Tree

```
In [25]: ► import pandas as pd
import os

def MyModel(X):
    model = tree.DecisionTreeClassifier(
        max_depth=10, # Limit tree depth to avoid excessive complexity
        min_samples_split=10, # Require at least 10 samples to split a node
        min_samples_leaf=5, # Require at least 5 samples per leaf node
        random_state=42)
    model = model.fit(X_train, y_train)

    predictions = model.predict(X)
    return predictions
```

```
In [26]: ► #random_out=holdout_data.iloc[:, :-1]
#X_test=random_out.sample(n=482,axis='rows')
```

```
In [27]: ▶ name = "My best model---Decision tree model: "
y_test_pred = MyModel(X_test)

# Evaluate model performance
accuracy_test = accuracy_score(y_test, y_test_pred)
recall_test = recall_score(y_test, y_test_pred, pos_label=1, average="binary")
precision_test = precision_score(y_test, y_test_pred, pos_label=1,
                                average="binary")

print("\t\t{m:s} Accuracy: {a:3.1%}, Recall {r:3.1%}, Precision {p:3.1%}".format
      (m=name,
       a=accuracy_test,
       r=recall_test,
       p=precision_test)
      )
```

```
My best model---Decision tree model: Accuracy: 84.1%, Recall 87.4%, Precision 81.8%
```

1. Initialize a Decision Tree classifier with:

- max_depth=10 (prevent overfitting)
- min_samples_split=10 (ensure splits happen on sufficiently large nodes)
- min_samples_leaf=5 (ensure leaf nodes are not too small)

2. Train the decision tree model on the training dataset (X_train, y_train).

3. Predict bankruptcy status for given input data (X).

4. Compute and return predictions.

5. Test the model using X_test and evaluate:

- Accuracy (percentage of correct predictions)
- Recall (how well bankrupt firms are identified)
- Precision (correctness of bankruptcy predictions)

6. Print model performance results.

For ECM: Easy to visualize decision rules, help to find which financial metrics drive bankruptcy risk

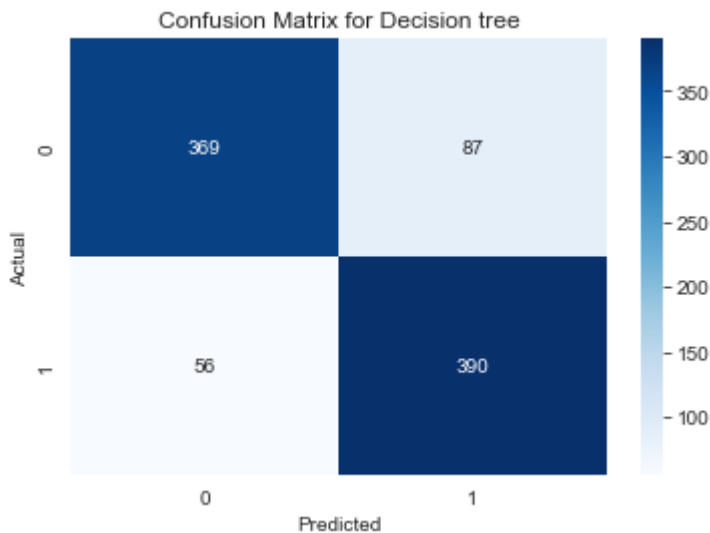
```
In [28]: ▶ cm = confusion_matrix(y_test, y_test_pred)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Decision tree')
plt.show()
```

Out[28]: <AxesSubplot:>

Out[28]: Text(0.5, 16.0, 'Predicted')

Out[28]: Text(34.0, 0.5, 'Actual')

Out[28]: Text(0.5, 1.0, 'Confusion Matrix for Decision tree')



I choose decision tree as my best model as its recall and precision are more in line with my expectation. As we know that "You may assume that it is 5 times worse to fail to identify a company that will go bankrupt than it is to fail to identify a company that won't go bankrupt", it indicates that we should have the recall higher than precision when trying to make trade-off between Recall and Precision. I set the test size to be small in order to achieve this goal which larger proportion of the data will be used for training the model. Thus the model has more training data to learn from and can better capture the patterns in the data that are associated with the positive class. However, SVM model and neural network model fail to achieve this goal. I try to increase the penalty parameter of the error term of SVM model "c", but even I increase to 20000 recall is still smaller than precision.

Decision Tree

Advantages:

- Easy to interpret and visualize.
- Captures non-linear relationships.
- Works well with small datasets and categorical variables.

Disadvantages:

- Prone to overfitting, especially with deep trees.
- Sensitive to noisy data.
- Not as effective for high-dimensional data without boosting or bagging.

Neural Network (MLP)

Advantages:

- Excellent for capturing complex, non-linear patterns.

- Works well with large datasets and high-dimensional data.
- Can model financial risk factors that interact in non-trivial ways.

Disadvantages:

- Hard to interpret (black-box model).
- Requires high computational power and tuning.
- Prone to overfitting if not properly regularized.

Logistic Regression

Advantages:

- Simple, interpretable, and easy to implement.
- Works well for linearly separable data.
- Computationally efficient, even for large datasets.

Disadvantages:

- Assumes linear relationships, which may not hold for complex financial data.
- Struggles with imbalanced data without resampling techniques.
- Less powerful compared to tree-based models or deep learning.

Support Vector Machine (SVM)

Advantages:

- Works well with small and medium-sized datasets.
- Effective for high-dimensional data.
- Can handle non-linear classification with the right kernel.

Disadvantages:

- Computationally expensive for large datasets.
- Hard to interpret compared to logistic regression or decision trees.
- Requires tuning of C and gamma to avoid overfitting or underfitting.

ECM:

Suggest for logistic regression (easy, interpretable) Neural network (For Complex Risk Patterns)

A single Decision Tree can easily memorize training data instead of learning general patterns (since variables for ECM interact in complex way)

SVM is quite expensive in large dataset